

# Feature and Circuit Identification through Sparse Eigendecomposition

---

## Some terms

### Model

The output of model  $f$  (of a given architecture), depends on two inputs - a set of parameters  $W$  and a set of feature  $x$ .

After training a given model, we end up with a set of weights/parameters  $W_0$ .

### Divergence

The divergence metric (not called loss, because I use that term later) should describe how much a model's output changes (for a given sample) when its parameters are altered.

In the case of a transformer, we can compute the normalized KL-divergence. The KL-divergence describes the divergence between probability distributions and we can normalize this metric to have a minimum of zero using the following equation (note that the unnormalized KL divergence does not necessarily have a minimum at 0). We could potentially use a normalized cross entropy here as well.

$$D(f, x, W) = \text{KL}(f(W, x), f(W_0, x)) - \text{KL}(f(W_0, x), f(W_0, x))$$

In the case of a regression problem, we can use the mean squared error loss as our divergence metric. We don't need to normalize it because MSE does have a guaranteed minimum at 0.

$$D(f, x, W) = \text{MSE}(f(W, x), f(W_0, x))$$

### Sample-level Hessian

We define the sample-level hessian as the second derivative of a model's loss with respect a set of parameters  $w$ , evaluated at  $W_0$  and  $X$ .

$$H(x, w) = \nabla^2_{\{w\}} D(f, x, W)$$

Using the Hessian, we can compute the second derivative of the loss with respect to any direction in weight/parameter space (given by vector  $u$ ) by:

$$\nabla^2_{\{u\}} D(f, x, W) = u(u^T H(x, w))$$

### Learning feature vectors

Our goal is to identify important directions ( $u$  or  $U$ ) in parameter space of this Hessian where: (1) For some samples,  $\nabla^2_{\{u\}} D$  is very high. *Moving parameters in the direction of  $u$  dramatically changes the output of the model.* (2) *There is low sample-level interference between each  $u$  vector.  $u$  vectors should either be nearly orthogonal, or if they*

have high cosine similarities, they should not both have high  $\|\nabla^2_{\mathbf{u}} D\|$  for a given sample.

We wish to learn a set of  $\mathbf{u}$ -vectors ( $\mathbf{U}$ ) that satisfy these conditions. We will therefore minimize two losses.

To satisfy the first goal we minimize:

$$L_{\text{low interference}}(\mathbf{x}) = \sum_{\mathbf{u}} (\|\nabla^2_{\mathbf{u}} D(f, \mathbf{x}, \mathbf{W}_0)\|)^2$$

To satisfy the second goal, we also want to minimize:

$$L_{\text{low interference}}(\mathbf{x}) = \|\nabla^2_{\mathbf{u}_1} D\| \text{space } \mathbf{u}_1 \|\nabla^2_{\mathbf{u}_2} D\| \text{space } \mathbf{u}_2^T |_{\mathbf{u}_1, \mathbf{u}_2}$$

## Optimization tricks

Optimizing losses that depend on a full hessian is computationally expensive. We use two tricks to minimize these losses.

### Jacobian-vector products

First, Instead of forming the full hessian, or full jacobian, we can compute nested jacobian-vector products.

$$\|\nabla^2_{\mathbf{u}} D(f, \mathbf{x}, \mathbf{W}_0)\| = \nabla_{\mathbf{w}} \left( \nabla_{\mathbf{w}} D(f, \mathbf{x}, \mathbf{W}_0) \text{space } \mathbf{u} \right) \text{space } \mathbf{u}$$

### Bi-level optimization

We use two separate optimization loops and perform bi-level optimization.

```

For x_batch in X:
    L1 = 0
    d2D_all = []

    # First optimization step.

    For u_batch in U:

        # Compute steep hessian loss.
        1. d2D(U) = second derivative of D(f, x, W_0) with respect to
u_batch.
        2. L1 = L1 + sum(d2D^2) # Compute first loss.
        3. d2D_all = d2D_all.append(d2D(U))

    U = U + L1*step    # Update U

    # Second optimization step.
    L2 = || (U d2D_all) (U d2D_all)^T || # Compute low-interference loss.
    U = U + (L1 + \lambda L2)*step # Update U.

```

Note that in our first optimization loop (minimizing  $L_{\text{low interference}}$ ), we can compute  $L$  in batches of  $U$ .

## Toy models

### XOR Model

We train a VERY simple neural network to learn the "XOR" function. The NN network consists of a single hidden layer of 2 nodes, with Gelu activation functions. The output of the hidden layer is summed to get a final output. There are 4 parameters today (2 weights and 2 biases). The training data for this network looks like:

```
[0, 1] --> 1
[1, 0] --> 1
[0, 0] --> 0
[1, 1] --> 0
```

### Toy model's of superposition

We train a TMS (autoencoder 5 features, 2 hidden dimensions, Relu activation, with  $W_{\text{in}}$  and  $W_{\text{out}}$  as transposes). We set the features to uniform random numbers between 0 and 1, with 5% sparsity. The TMS model successfully represents the features in pentagonal superposition.

### Transformers

We use the tiny-stories-1M transformer.

## Results

### XOR Model

The eigenmodel successfully finds features that are most highly activated by [0,1], [1,1], and [1,0].

```
feature_idx
[sample input values] -> feature value

feature 0
[1. 1.] -> 2.3252432
[1. 0.] -> 1.0929958
[0. 1.] -> 0.61902356
[0. 0.] -> 0.273664

feature 1
[0. 1.] -> 2.670694
[1. 0.] -> 0.65577537
[1. 1.] -> 0.52892005
[0. 0.] -> 0.15629882
```

```
feature 2
[1. 0.] -> 2.021866
[1. 1.] -> 0.39354768
[0. 1.] -> 0.035297774
[0. 0.] -> 0.011965705
```

## TMS

The most highly activating samples are the following:

```
feature_idx
[sample input values] -> feature value

feature 0
[0.    0.779 0.    0.    0.948] -> 3.342
[0.    0.833 0.    0.    0.35 ] -> 2.111
[0.    0.999 0.    0.    0.   ] -> 1.862
[0.    0.997 0.    0.    0.   ] -> 1.857
[0.    0.272 0.    0.    0.944] -> 1.836

feature 1
[0.888 0.    0.    0.835 0.   ] -> 3.406
[0.795 0.    0.    0.812 0.   ] -> 3.075
[0.796 0.    0.    0.78  0.   ] -> 2.977
[0.472 0.    0.    0.881 0.   ] -> 2.455
[0.285 0.    0.    0.91  0.   ] -> 2.101

feature 2
[0.758 0.    0.    0.    0.975] -> 2.078
[0.641 0.    0.    0.    0.698] -> 1.383
[0.517 0.    0.    0.    0.744] -> 1.295
[0.    0.    0.    0.    0.999] -> 1.274
[0.    0.    0.    0.    0.992] -> 1.259

feature 3
[0.758 0.    0.    0.    0.975] -> 2.516
[0.641 0.    0.    0.    0.698] -> 1.724
[0.979 0.    0.    0.    0.   ] -> 1.666
[0.971 0.    0.    0.    0.   ] -> 1.646
[0.517 0.    0.    0.    0.744] -> 1.546

feature 4
[0.126 0.    0.942 0.477 0.   ] -> 2.439
[0.    0.    0.984 0.    0.   ] -> 2.241
[0.    0.    0.971 0.    0.   ] -> 2.199
[0.    0.    0.97 0.    0.   ] -> 2.196
[0.    0.    0.967 0.    0.   ] -> 2.187
```

If we only consider completely sparse samples, the top features are the following.

```
feature_idx
[sample input values] -> feature value
```

```
feature 0
[0.    0.999 0.    0.    0.    ] -> 1.862
[0.    0.997 0.    0.    0.    ] -> 1.857
[0.    0.981 0.    0.    0.    ] -> 1.816
[0.    0.971 0.    0.    0.    ] -> 1.793
[0.    0.957 0.    0.    0.    ] -> 1.758
```

```
feature 1
[0.    0.    0.    0.998 0.    ] -> 1.728
[0.    0.    0.    0.996 0.    ] -> 1.724
[0.    0.    0.    0.995 0.    ] -> 1.721
[0.    0.    0.    0.991 0.    ] -> 1.712
[0.    0.    0.    0.97 0.    ] -> 1.662
```

```
feature 2
[0.    0.    0.    0.    0.999] -> 1.274
[0.    0.    0.    0.    0.992] -> 1.259
[0.    0.    0.    0.    0.969] -> 1.213
[0.    0.    0.    0.    0.952] -> 1.181
[0.    0.    0.    0.    0.938] -> 1.154
```

```
feature 3
[0.979 0.    0.    0.    0.    ] -> 1.666
[0.971 0.    0.    0.    0.    ] -> 1.646
[0.915 0.    0.    0.    0.    ] -> 1.523
[0.901 0.    0.    0.    0.    ] -> 1.493
[0.897 0.    0.    0.    0.    ] -> 1.485
```

```
feature 4
[0.    0.    0.984 0.    0.    ] -> 2.241
[0.    0.    0.971 0.    0.    ] -> 2.199
[0.    0.    0.97 0.    0.    ] -> 2.196
[0.    0.    0.967 0.    0.    ] -> 2.187
[0.    0.    0.958 0.    0.    ] -> 2.157
```

## Transformer

Here are some results using just the weights in `transformer.blocks.4.attn.W_K` (4096 weights) in the `tiny-stories-1M` model. Here are some results from a model trained very briefly (<10 epochs, on only 500 token sets) with only 10 features pulled out. Results on data the eigenmodel was not trained on.

```
Feature description (by me)
tokens (activating token bolded) -> Feature value
```

```
Word/punctuation after a name
upon a time, there was a kind man named Tom**. ** Tom had a big -> .
```

(Value: 42.845)  
 zoo.Once upon a time, there was a little boy named Timmy\*\*. \*\* -> . (Value: 40.725)  
 upon a time, there was a woman named Lily. \*\* She\*\* loved to go for -> She (Value: 38.833)  
 Once upon a time, there was a little boy named Timmy\*\*. \*\* Timmy -> . (Value: 38.542)  
 Once upon a time, there was a little boy named Timmy\*\*. \*\* Timmy -> . (Value: 38.542)

Once upon a time  
 upon a time, there was a woman named Lily. \*\* She\*\* loved to go for -> She (Value: 64.518)  
 could always ask the black cat.Once upon a time\*\*, \*\* there was a small -> , (Value: 56.562)  
 loved ones. The end.Once upon a time\*\*, \*\* there was a little girl -> , (Value: 50.217)  
 enjoyed the sunshine.Once upon a time\*\*, \*\* there was a little boy -> , (Value: 46.414)  
 street and always looked out for banana peels.Once upon a time\*\*, \*\* there -> , (Value: 46.287)

Tim  
 and trucks. One day, \*\* Tim\*\*my's dad took him to the park to -> Tim (Value: 69.034)  
 lost and felt sad.newlinenewlineOne day, \*\* Tim\*\*my's friend, a -> Tim (Value: 66.466)  
 toys. The next day, \*\* Tim\*\*my went to his friend's house and said -> Tim (Value: 62.213)  
 was playing just as well as before. \*\* Tim\*\*my was so happy with his new -> Tim (Value: 57.687)  
 the noisy trunk.Once upon a time, there was a boy named\*\* Tim\*\*my -> Tim (Value: 54.684)

?  
 park. One day, she saw a big statue of a dog\*\*. \*\* It was -> . (Value: 30.722)  
 't like that. He wanted his truck to be the\*\* best\*\*. They argued for -> best (Value: 27.736)  
 every day. One day, he saw a little bird on a branch\*\*. \*\* The -> . (Value: 26.222)  
 newlinenewlineOne day, Timmy saw a black cat in his backyard\*\*. \*\* The -> . (Value: 22.828)  
 ily said, "Let's ask for help!" They saw a man\*\* and\*\* asked -> and (Value: 22.407)

Gender / polysemanic  
 upon a time, there was a woman named Lily. \*\* She\*\* loved to go for -> She (Value: 59.775)  
 Once upon a time, there was an old man. \*\* He\*\* liked to read magazines -> He (Value: 39.421)  
 the noisy trunk.Once upon a time, there was a boy named\*\* Tim\*\*my -> Tim (Value: 38.685)  
 summer.Once upon a time, there was an old lady. \*\* She\*\* was very -> She

(Value: 37.195)  
came in to see what was going on.newlinenewlineShe told\*\* Tim\*\*my that ->  
Tim (Value: 36.652)

saw + polysemantic  
the park every day. newlinenewlineOne day, the man\*\* saw\*\* the woman ->  
saw (Value: 39.632)  
the dark hole in the ground and he fell in.newlinenewline\*\*Tim\*\*my tried -  
> Tim (Value: 30.418)  
and went inside the tent. newlinenewlineThe lion\*\* saw\*\* a man with a ->  
saw (Value: 28.665)  
Allowed." Timmy was sad because he wanted Max to come with them\*\*.\*\* -> .  
(Value: 28.455)  
's mommy lifted her up so she could see over it. She\*\* saw\*\* her -> saw  
(Value: 28.105)

Tim + "then on?"  
toy tools. From that day on,\*\* Tim\*\*my learned the importance of sharing  
and -> Tim (Value: 75.689)  
was playing just as well as before.\*\* Tim\*\*my was so happy with his new ->  
Tim (Value: 63.231)  
at them. He even got to touch a baby shark! After that,\*\* Tim\*\* -> Tim  
(Value: 59.772)  
then, his mom came in and asked what was wrong.\*\* Tim\*\*my told her -> Tim  
(Value: 58.194)  
the dark hole in the ground and he fell in.newlinenewline\*\*Tim\*\*my tried -  
> Tim (Value: 57.360)

Animals?  
named Timmy. He had a big, brown dog named Max\*\*.\*\* Max loved -> . (Value:  
54.549)  
park. One day, she saw a big statue of a dog\*\*.\*\* It was -> . (Value:  
43.561)  
She loved to walk on the trail with her dog, Max\*\*.\*\* Max was very -> .  
(Value: 39.359)  
newlinenewlineOne day, Timmy saw a black cat in his backyard\*\*.\*\* The -> .  
(Value: 36.026)  
and wanted to make him feel better. Tommy\*\* told\*\* Sammy a joke and Sammy  
laughed -> told (Value: 34.926)

was playing just as well as before.\*\* Tim\*\*my was so happy with his new ->  
Tim (Value: 55.711)  
part of the park where dogs were allowed.\*\* Tim\*\*my was happy again  
because he -> Tim (Value: 38.506)  
then, his mom came in and asked what was wrong.\*\* Tim\*\*my told her -> Tim  
(Value: 37.558)  
should answer with courage. So,\*\* Tim\*\*my took a deep breath and stood up  
-> Tim (Value: 34.767)  
came in to see what was going on.newlinenewlineShe told\*\* Tim\*\*my that ->  
Tim (Value: 33.680)

```
came in to see what was going on.newlinenewlineShe told** Tim**my that ->
Tim (Value: 46.779)
at them. He even got to touch a baby shark! After that,** Tim** -> Tim
(Value: 34.167)
the dark hole in the ground and he fell in.newlinenewline**Tim**my tried -
> Tim (Value: 30.990)
part of the park where dogs were allowed.** Tim**my was happy again
because he -> Tim (Value: 30.866)
that they were safe.newlinenewlineBut the next day,** Tim**my didn't ->
Tim (Value: 30.029)
```